

# Data analysis Project on Energy Consumption Dataset

## Libraries and Configurations

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import kagglehub

sns.set_theme(style="whitegrid")
plt.rcParams['figure.figsize'] = (12, 6)
```

## Data loading

```
path = kagglehub.dataset_download("mrsimple07/energy-consumption-prediction")
print("Path to dataset files:", path)
Path to dataset files: /Users/nameisalfio/.cache/kagglehub/datasets/mrsimple07/energy-consumption-prediction
df = pd.read_csv(path + "/Energy_consumption.csv")

print(df.info())
print(df.describe())

<class 'pandas.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Timestamp        1000 non-null    str    
 1   Temperature     1000 non-null    float64 
 2   Humidity         1000 non-null    float64 
 3   SquareFootage   1000 non-null    float64 
 4   Occupancy        1000 non-null    int64  
 5   HVACUsage        1000 non-null    str    
 6   LightingUsage   1000 non-null    str    
 7   RenewableEnergy  1000 non-null    float64 
 8   DayOfWeek       1000 non-null    str    
 9   Holiday          1000 non-null    str    
 10  EnergyConsumption 1000 non-null    float64 
dtypes: float64(5), int64(1), str(5)
memory usage: 86.1 KB
None
Temperature      Humidity      SquareFootage    Occupancy    RenewableEnergy \
count  1000.000000  1000.000000  1000.000000  1000.000000  1000.000000
```

mean	24.982026	45.395412	1500.052488	4.581000	15.132813
std	2.836850	8.518905	288.418873	2.865598	8.745917
min	20.007565	30.015975	1000.512661	0.000000	0.006642
25%	22.645070	38.297722	1247.108548	2.000000	7.628385
50%	24.751637	45.972116	1507.967426	5.000000	15.072296
75%	27.418174	52.420066	1740.340165	7.000000	22.884064
max	29.998671	59.969085	1999.982252	9.000000	29.965327

EnergyConsumption	
count	1000.000000
mean	77.055873
std	8.144112
min	53.263278
25%	71.544690
50%	76.943696
75%	82.921742
max	99.201120

## Data Preprocessing & Time Engineering

We need to transform the timestamp into a machine-readable format and handle categorical variables.

```
print(f'Missing values for each column: {df.isna().sum()}')

Missing values for each column: Timestamp          0
Temperature        0
Humidity          0
SquareFootage     0
Occupancy         0
HVACUsage         0
LightingUsage     0
RenewableEnergy    0
DayOfWeek         0
Holiday           0
EnergyConsumption 0
dtype: int64

# 1. Convert Timestamp to datetime objects
df['Timestamp'] = pd.to_datetime(df['Timestamp'])

# 2. Extract temporal features
# These are essential for identifying daily and weekly patterns
df['Hour'] = df['Timestamp'].dt.hour
df['Day'] = df['Timestamp'].dt.day
df['Month'] = df['Timestamp'].dt.month
df['DayOfWeek_Num'] = df['Timestamp'].dt.dayofweek # 0=Monday, 6=Sunday
```

```

# 3. Check for missing values explicitly
missing_data = df.isnull().sum()
if missing_data.sum() == 0:
    print("Data Integrity Check: No missing values found.")
else:
    print(f"Missing values found:\n{missing_data[missing_data > 0]}")

# 4. Brief look at the updated dataframe
print("\nFirst 5 rows with new temporal features:")
print(df[['Timestamp', 'Hour', 'DayOfWeek', 'EnergyConsumption']].head())
Data Integrity Check: No missing values found.

First 5 rows with new temporal features:
   Timestamp  Hour  DayOfWeek  EnergyConsumption
0 2022-01-01  00:00:00    0     Monday        75.364373
1 2022-01-01  01:00:00    1   Saturday        83.401855
2 2022-01-01  02:00:00    2   Sunday         78.270888
3 2022-01-01  03:00:00    3 Wednesday        56.519850
4 2022-01-01  04:00:00    4   Friday         70.811732

```

### Univariate Analysis (Target Variable)

Before looking for relationships, we need to understand how energy consumption “moves” on its own. We look for outliers that could compromise future models.

```

# Setup the visualization grid
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Histogram & KDE: Check for normality
sns.histplot(df['EnergyConsumption'], kde=True, ax=axes[0], color='skyblue')
axes[0].set_title('Energy Consumption Distribution')
axes[0].set_xlabel('Consumption (kWh)')

# Boxplot: Check for statistical outliers
sns.boxplot(x=df['EnergyConsumption'], ax=axes[1], color='lightcoral')
axes[1].set_title('Energy Consumption Boxplot')
axes[1].set_xlabel('Consumption (kWh)')

plt.tight_layout()
plt.show()

print(f"Skewness: {df['EnergyConsumption'].skew():.2f}") # should be near 0, indicating symmetry
print(f"Kurtosis (Fisher): {df['EnergyConsumption'].kurt():.2f}") # should be near 0, indicating normality

```

Skewness: 0.03  
Kurtosis (Fisher): -0.30




Figure 1: png

```

# Quantitative outlier detection using IQR
Q1 = df['EnergyConsumption'].quantile(0.25)
Q2 = df['EnergyConsumption'].quantile(0.5)      # Median
Q3 = df['EnergyConsumption'].quantile(0.75)
Q4 = df['EnergyConsumption'].quantile(1.0)      # Max
mean = df['EnergyConsumption'].mean()
dev_std = df['EnergyConsumption'].std()

IQR = Q3 - Q1      # Interquartile Range
outlier_threshold = 1.5 * IQR

outliers = df[(df['EnergyConsumption'] < Q1 - outlier_threshold) |
              (df['EnergyConsumption'] > Q3 + outlier_threshold)]

print("Statistical Summary for Target Variable:")
print("-" * 50)
print(f"First quantile (Q1): {Q1:.2f}")
print(f"Median (Q2): {Q2:.2f}")
print(f"Third quantile (Q3): {Q3:.2f}")
print(f"Maximum (Q4): {Q4:.2f}")
print(f"Mean: {mean:.2f}")
print(f"Standard deviation: {dev_std:.2f}\t Variance: {dev_std**2:.2f}")
print(f"\nPotential outliers detected: {len(outliers)}\n{outliers}")

Statistical Summary for Target Variable:
-----
First quantile (Q1): 71.54
Median (Q2): 76.94
Third quantile (Q3): 82.92
Maximum (Q4): 99.20
Mean: 77.06
Standard deviation: 8.14      Variance: 66.33

```

```
Potential outliers detected: 1:
```

```
    Timestamp  Temperature  Humidity  SquareFootage  Occupancy  \
69  2022-01-03 21:00:00      20.735716  48.506636      1836.542651      4
                                               HVACUsage  LightingUsage  RenewableEnergy  DayOfWeek  Holiday  \
69          Off                  On              9.295439     Friday        No
                                               EnergyConsumption  Hour  Day  Month  DayOfWeek_Num
69            53.263278       21    3     1                   0
```

### Multivariate Analysis (The “Drivers”)

Now let's identify the real drivers of consumption. We will create a correlation matrix to see which numerical variables (Temperature, Humidity, Employment) are most closely related to energy.

```
# Select only numerical columns for correlation
numerical_df = df.select_dtypes(include=[np.number])

# Compute correlation matrix
corr_matrix = numerical_df.corr()

# Plot heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='RdYlGn', center=0, fmt='.2f')
plt.title('Feature Correlation Heatmap: What drives consumption?')
plt.show()

mask = np.triu(np.ones_like(corr_matrix, dtype=bool)) #exclude the upper triangular matrix
corr_lower = corr_matrix.mask(mask)

corr_pairs = corr_lower.stack() #convert in tuples
strong_corr = corr_pairs[abs(corr_pairs) >= 0.5].sort_values(key=np.abs, ascending=False)
print("Strong correlations:")
print("-" * 50)
print(strong_corr)

Strong correlations:
-----
EnergyConsumption  Temperature      0.69641
dtype: float64
```

### Bivariate Analysis (Deep Dive into Drivers)

We will now visualize the relationship between Temperature and Consumption, and then analyze how categorical variables like HVACUsage and Occupancy



Figure 2: png

shift the consumption baseline.

```
# Visualize the linear relationship between Temperature and Energy Consumption
plt.figure(figsize=(10, 6))
sns.replot(data=df, x='Temperature', y='EnergyConsumption',
            scatter_kws={'alpha':0.3}, line_kws={'color':'red'})
plt.title('Bivariate Analysis: Energy Consumption vs Temperature')
plt.xlabel('Temperature')
plt.ylabel('Energy Consumption (kWh)')
plt.show()
```




Figure 3: png

```
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Boxplot for Occupancy
sns.boxplot(
    data=df,
    x='Occupancy',
    y='EnergyConsumption',
    hue='Occupancy',
    palette='Set2',
    ax=axes[0],
    dodge=False,
    legend=False
)
axes[0].set_title('Energy Consumption by Occupancy Level')
```

```

# Boxplot for HVACUsage
sns.boxplot(
    data=df,
    x='HVACUsage',
    y='EnergyConsumption',
    hue='HVACUsage',
    palette='Set3',
    ax=axes[1],
    dodge=False,
    legend=False
)
axes[1].set_title('Energy Consumption by HVAC Status')

plt.tight_layout()
plt.show()

```




Figure 4: png

- **Occupancy Paradox** : Nota come il consumo mediano per Occupancy = 0 sia molto vicino a quello per Occupancy = 9. Questo è un segnale d'allarme critico: l'edificio consuma quasi la stessa energia sia quando è vuoto che quando è pieno.
- **HVAC Shift** : Il grafico a destra conferma che lo stato “On” del sistema HVAC(Heating, Ventilation, and Air Conditioning) alza l'intero intervallo di consumo, ma esiste un consumo significativo anche quando il sistema è “Off”.

### Hourly Load Profile Analysis

```

# Create an average hourly profile to identify peak/off-peak patterns
hourly_summary = df.groupby('Hour')[['EnergyConsumption']].agg(['mean', 'std']).reset_index()

plt.figure(figsize=(12, 6))

```

```

plt.plot(hourly_summary['Hour'], hourly_summary['mean'], color='darkblue', marker='o', linewidth=2)
plt.fill_between(hourly_summary['Hour'],
                 hourly_summary['mean'] - hourly_summary['std'],
                 hourly_summary['mean'] + hourly_summary['std'],
                 color='blue', alpha=0.1, label='Standard Deviation')

plt.title('Average Daily Load Profile (Hourly Analysis)', fontsize=14)
plt.xlabel('Hour of the Day', fontsize=12)
plt.ylabel('Energy Consumption (kWh)', fontsize=12)
plt.xticks(range(0, 24))
plt.legend()
plt.grid(True, linestyle='--', alpha=0.6)
plt.show()

```




Figure 5: png

- **Time-series Insight :** Look for peaks during working hours. A flat line indicates inefficient constant usage.

### Baseload Calculation

```

# Calculate the empirical Baseload (10th percentile to avoid outliers)
baseload = df['EnergyConsumption'].quantile(0.1)
peak_load = df['EnergyConsumption'].max()
dynamic_ratio = (peak_load - baseload) / peak_load

print(f"--- Energy Efficiency Metrics ---")
print(f"Empirical Baseload: {baseload:.2f} kWh")
print(f"Peak Demand: {peak_load:.2f} kWh")

```

```

print(f"Dynamic Ratio: {dynamic_ratio:.2%}")

--- Energy Efficiency Metrics ---
Empirical Baseload: 66.46 kWh
Peak Demand: 99.20 kWh
Dynamic Ratio: 33.00%

```

- **Logic** : High baseload (>60% of peak) usually suggests poorly configured systems or 24/7 waste.

### Renewable Energy Impact & Net Consumption

We need to understand whether renewable energy production occurs when it is most needed or whether it is wasted.

```

# 1. Calculate Net Consumption (Grid reliance)
df['Net_Consumption'] = df['EnergyConsumption'] - df['RenewableEnergy']

# 2. Visualize the offset
plt.figure(figsize=(12, 6))
sns.lineplot(data=df.head(100), x='Timestamp', y='EnergyConsumption', label='Gross Consumption')
sns.lineplot(data=df.head(100), x='Timestamp', y='RenewableEnergy', label='Renewable Production')
plt.title('Gross Consumption vs Renewable Production (First 100 samples)')
plt.ylabel('kWh')
plt.show()

print(f"Average Renewable Contribution: {((df['RenewableEnergy'].mean() / df['EnergyConsumption'].mean()) * 100):.2f} %")

```




Figure 6: png

Average Renewable Contribution: 19.64%

## Energy Use Intensity (EUI)

In power engineering, absolute consumption doesn't say it all. A 100m<sup>2</sup> building that consumes 80kWh is inefficient compared to a 1000m<sup>2</sup> building that consumes the same.

```
df['EUI'] = df['EnergyConsumption'] / df['SquareFootage']
print(f"Mean EUI: {df['EUI'].mean():.4f} kWh/sqft")
Mean EUI: 0.0534 kWh/sqft
```

## Holiday and Weekend vs Weekday Analysis

Buildings have different “behaviors” on holidays. If consumption during a Holiday is identical to a working Tuesday, there is a serious error in the system planning (e.g. lights and heating turned on unnecessarily).

```
holiday_profile = df[df['Holiday'] == 'Yes'].groupby('Hour')['EnergyConsumption'].mean()
workday_profile = df[df['Holiday'] == 'No'].groupby('Hour')['EnergyConsumption'].mean()

plt.figure(figsize=(12, 6))
plt.plot(workday_profile, label='Workday', color='blue')
plt.plot(holiday_profile, label='Holiday', color='red', linestyle='--')
plt.legend()
plt.title('Hourly Profile: Workdays vs Holidays')
plt.show()
```




Figure 7: png

In a properly managed building, we would expect the blue line (Workday) to be consistently above the red line (Holiday), especially during working hours. Instead, let's observe the opposite:

- Higher consumption on holidays: The red line shows consumption peaks (especially around 1, 10 and 19 am) that significantly exceed working day levels.
- Management Inconsistency: This suggests that systems (HVAC, lights) are not turned off or reduced during holidays. On the contrary, they appear to operate under higher loads, which is a clear sign of massive energy waste or incorrect timer scheduling.

## Statistical Significance and Feature Selection (Backward Elimination)

```

import statsmodels.api as sm
from sklearn.preprocessing import LabelEncoder

# 1. Encoding categorical variables for statistical analysis
le = LabelEncoder()
df_stat = df.copy()

categorical_to_encode = ['HVACUsage', 'LightingUsage', 'Holiday']
for col in categorical_to_encode:
    df_stat[col] = le.fit_transform(df_stat[col].astype(str))

# 2. Define initial features (excluding non-numerical and target variables)
features = [
    'Temperature', 'Humidity', 'SquareFootage', 'Occupancy',
    'HVACUsage', 'LightingUsage', 'RenewableEnergy',
    'Hour', 'DayOfWeek_Num', 'Holiday'
]

X = df_stat[features]
y = df_stat['EnergyConsumption']

# 3. Add a constant (intercept) as required by statsmodels for OLS
X = sm.add_constant(X)

# 4. Backward Elimination Process
def perform_backward_elimination(X, y, alpha=0.05):
    current_features = X.columns.tolist()
    while len(current_features) > 0:
        # Fit the OLS model
        model = sm.OLS(y, X[current_features]).fit()

        # Get the feature with the highest p-value
        max_p_value = model.pvalues.max()
        if max_p_value > alpha:

```

```

        excluded_feature = model.pvalues.idxmax()
        print(f"Removing '{excluded_feature}' (p-value: {max_p_value:.4f})")
        current_features.remove(excluded_feature)
    else:
        break
    return model

# 5. Execute and Display Results
final_stat_model = perform_backward_elimination(X, y)

print("\n" + "="*20 + " FINAL STATISTICAL SUMMARY " + "="*20)
print(final_stat_model.summary())

Removing 'SquareFootage' (p-value: 0.5861)
Removing 'DayOfWeek_Num' (p-value: 0.3839)
Removing 'Hour' (p-value: 0.3239)
Removing 'Holiday' (p-value: 0.3058)

=====
          FINAL STATISTICAL SUMMARY
          OLS Regression Results
=====
Dep. Variable:      EnergyConsumption    R-squared:                 0.619
Model:                  OLS    Adj. R-squared:              0.617
Method:                Least Squares   F-statistic:                 269.4
Date: Mon, 26 Jan 2026   Prob (F-statistic):        2.22e-204
Time:           22:10:21    Log-Likelihood:             -3032.6
No. Observations:      1000    AIC:                      6079.
Df Residuals:         993    BIC:                      6114.
Df Model:                   6
Covariance Type:     nonrobust
=====

            coef    std err        t      P>|t|      [0.025      0.975]
-----
const      22.4524     1.741     12.896     0.000     19.036     25.869
Temperature  1.9935     0.056     35.450     0.000     1.883     2.104
Humidity    -0.0402     0.019     -2.136     0.033     -0.077     -0.003
Occupancy     0.5281     0.056      9.489     0.000      0.419     0.637
HVACUsage     4.6127     0.321     14.385     0.000      3.983     5.242
LightingUsage  1.7050     0.319      5.338     0.000      1.078     2.332
RenewableEnergy  0.0727     0.018      3.988     0.000      0.037     0.109
=====

Omnibus:                 0.071 Durbin-Watson:                 2.023
Prob(Omnibus):            0.965 Jarque-Bera (JB):               0.126
Skew:                     -0.012 Prob(JB):                  0.939
Kurtosis:                  2.950 Cond. No.                    598.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

### Visualization of Significant Drivers

Once the non-significant features are removed, we can plot the coefficients to see which variables have the greatest weight (positive or negative) on energy consumption.

```
# Extract coefficients and p-values from the final model
coef_df = pd.DataFrame({
    'Feature': final_stat_model.params.index,
    'Coefficient': final_stat_model.params.values,
    'P-Value': final_stat_model.pvalues.values
}).sort_values(by='Coefficient', ascending=False)

coef_df = coef_df[coef_df['Feature'] != 'const']

plt.figure(figsize=(10, 6))
sns.barplot(
    data=coef_df,
    x='Coefficient',
    y='Feature',
    hue='Feature',
    palette='viridis',
    legend=False
)

plt.title('Impact of Features on Energy Consumption (Final OLS Model)')
plt.axvline(0, color='black', linestyle='--', linewidth=1)
plt.show()
```

### Verification of Residuals

A key assumption for the validity of the t-test and F-test is that residuals are normally distributed and independent.

```
# Plotting Residuals Distribution
residuals = final_stat_model.resid

fig, axes = plt.subplots(1, 2, figsize=(16, 5))

sns.histplot(residuals, kde=True, ax=axes[0], color='gray')
axes[0].set_title('Residuals Distribution (Normality Check)')

sns.scatterplot(x=final_stat_model.fittedvalues, y=residuals, ax=axes[1], alpha=0.5)
```




Figure 8: png

```

axes[1].axhline(0, color='red', linestyle='--')
axes[1].set_title('Residuals vs Fitted Values (Homoscedasticity Check)')
axes[1].set_xlabel('Predicted Consumption')
axes[1].set_ylabel('Residuals')

plt.show()

```




Figure 9: png

```

selected_features = final_stat_model.pvalues[final_stat_model.pvalues < 0.05].index.tolist()
if 'const' in selected_features: selected_features.remove('const')

print("Selected features for modeling based on statistical significance:")
print(selected_features)

```

```
Selected features for modeling based on statistical significance:  
['Temperature', 'Humidity', 'Occupancy', 'HVACUsage', 'LightingUsage', 'RenewableEnergy']
```

## Predictive Modelling

---

### Data preprocessing

```
import torch  
import torch.nn as nn  
import torch.optim as optim  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler, LabelEncoder  
  
# 1. Feature filtering based on statistical significance  
# This list is derived from the p-values of the preliminary statistical model  
selected_features = final_stat_model.pvalues[final_stat_model.pvalues < 0.05].index.tolist()  
if 'const' in selected_features:  
    selected_features.remove('const')  
  
# 2. Data Preparation and Categorical Encoding  
le = LabelEncoder()  
df_ml = df.copy()  
  
# Identify which categorical columns are present in the selected features  
categorical_targets = ['HVACUsage', 'LightingUsage', 'DayOfWeek', 'Holiday']  
active_categoricals = [col for col in categorical_targets if col in selected_features]  
  
for col in active_categoricals:  
    df_ml[col] = le.fit_transform(df_ml[col].astype(str))  
  
# 3. Dynamic Feature Selection  
# X is constructed using only the statistically significant predictors  
X = df_ml[selected_features].values  
y = df_ml['EnergyConsumption'].values.reshape(-1, 1)  
  
# 4. Dataset Partitioning (Train/Test Split)  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
# 5. Feature and Target Scaling  
# Standardizing features to mean=0 and variance=1 for optimal gradient descent  
scaler_X = StandardScaler()  
scaler_y = StandardScaler()  
  
X_train = scaler_X.fit_transform(X_train)
```

```

X_test = scaler_X.transform(X_test)
y_train = scaler_y.fit_transform(y_train)
y_test = scaler_y.transform(y_test)

# 6. Tensor Conversion for PyTorch Compatibility
X_train_t = torch.tensor(X_train, dtype=torch.float32)
y_train_t = torch.tensor(y_train, dtype=torch.float32)
X_test_t = torch.tensor(X_test, dtype=torch.float32)
y_test_t = torch.tensor(y_test, dtype=torch.float32)

print(f"Training set size: {X_train_t.shape[0]} samples with {X_train_t.shape[1]} features.")
print(f"Testing set size: {X_test_t.shape[0]} samples.")

Training set size: 800 samples with 6 features.
Testing set size: 200 samples.

```

### Definition of Regression model

```

class EnergyRegressor(nn.Module):
    def __init__(self, input_dim):
        super(EnergyRegressor, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(input_dim, 64),
            nn.BatchNorm1d(64),
            nn.ReLU(),

            nn.Linear(64, 32),
            nn.BatchNorm1d(32),
            nn.ReLU(),

            nn.Linear(32, 1)
        )

    def forward(self, x):
        return self.net(x)

```

### Training Loop

We will use the MSE loss function defined as:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

```

from torch.utils.data import DataLoader, TensorDataset

train_ds = TensorDataset(X_train_t, y_train_t)
test_ds = TensorDataset(X_test_t, y_test_t)

```

```

train_loader = DataLoader(train_ds, batch_size=32, shuffle=True)
val_loader = DataLoader(test_ds, batch_size=len(test_ds))
print(f"Batches in training loader: {len(train_loader)}")

Batches in training loader: 25

model = EnergyRegressor(X_train_t.shape[1])
print(model)
print('-' * 50)

optimizer = optim.Adam(model.parameters(), lr=0.0001)
criterion = nn.MSELoss()

epochs = 50
train_losses = []
val_losses = []

for epoch in range(epochs):
    # --- Training phase ---
    model.train()
    running_train_loss = 0.0

    for batch_X, batch_y in train_loader:
        optimizer.zero_grad()

        # Forward pass on the mini-batch
        outputs = model(batch_X)
        loss = criterion(outputs, batch_y)

        # Backward pass
        loss.backward()
        optimizer.step()

        running_train_loss += loss.item()

    # Average training loss for the batch
    epoch_train_loss = running_train_loss / len(train_loader)
    train_losses.append(epoch_train_loss)

    # --- Validation phase ---
    model.eval()
    running_val_loss = 0.0

    with torch.no_grad():
        for val_X, val_y in val_loader:
            prediction = model(val_X)

```

```

        v_loss = criterion(prediction, val_y)
        running_val_loss += v_loss.item()

    epoch_val_loss = running_val_loss / len(val_loader)
    val_losses.append(epoch_val_loss)

    if (epoch + 1) % 50 == 0:
        print(f"Epoch [{epoch+1}/{epoch+1}] | Avg Train Loss: {epoch_train_loss:.4f} | Avg Val Loss: {epoch_val_loss:.4f}")

EnergyRegressor(
    (net): Sequential(
        (0): Linear(in_features=6, out_features=64, bias=True)
        (1): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): ReLU()
        (3): Linear(in_features=64, out_features=32, bias=True)
        (4): BatchNorm1d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (5): ReLU()
        (6): Linear(in_features=32, out_features=1, bias=True)
    )
)
-----
Epoch [50/50] | Avg Train Loss: 0.3806 | Avg Val Loss: 0.4111

plt.figure(figsize=(10, 6))
plt.plot(train_losses, label='Train Loss')
plt.plot(val_losses, label='Validation Loss')
plt.title('Training and Validation curves')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.legend()
plt.show()

from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

model.eval()

with torch.no_grad():
    predictions_scaled = model(X_test_t)
    predictions_numpy = predictions_scaled.cpu().numpy()
    actuals_numpy = y_test_t.cpu().numpy()

    # Inverse transform per tornare alla scala reale kWh
    predictions_kwh = scaler_y.inverse_transform(predictions_numpy)
    actuals_kwh = scaler_y.inverse_transform(actuals_numpy)

    # Visualizing Predictions vs Actual (in kWh)
    plt.figure(figsize=(10, 6))

```




Figure 10: png

```

plt.scatter(actuals_kwh, predictions_kwh, alpha=0.5, color='teal', label='Regression Predict')
plt.plot([actuals_kwh.min(), actuals_kwh.max()], [actuals_kwh.min(), actuals_kwh.max()],
         'r--', lw=2, label='Perfect Fit')

plt.xlabel('Actual Consumption (kWh)')
plt.ylabel('Predicted Consumption (kWh)')
plt.title('Baseline Analysis: Actual vs Predicted Energy (Real Scale)')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.6)
plt.show()

mae_kwh = mean_absolute_error(actuals_kwh, predictions_kwh)
mse_kwh = mean_squared_error(actuals_kwh, predictions_kwh)
r2 = r2_score(actuals_kwh, predictions_kwh)

print(f"--- Final Metrics of the Regression Model ---")
print(f"Mean Absolute Error (MAE): {mae_kwh:.2f} kWh")
print(f"Mean Squared Error (MSE): {mse_kwh:.2f} kWh^2")
print(f"R2 Score: {r2:.4f}")

--- Final Metrics of the Regression Model ---
Mean Absolute Error (MAE): 4.12 kWh
Mean Squared Error (MSE): 27.32 kWh^2
R2 Score: 0.5829

```




Figure 11: png

```

from sklearn.metrics import r2_score, explained_variance_score

# 1. Calculation of the Coefficient of Determination (R2 Score)
# R2 represents the proportion of variance for the dependent variable
# that is explained by the independent variables in the model.
r2 = r2_score(actuals_kwh, predictions_kwh)

# 2. Calculation of the Explained Variance Score
# This score measures the proportion to which a mathematical model accounts
# for the variation (dispersion) of a given data set.
evs = explained_variance_score(actuals_kwh, predictions_kwh)

print("-" * 50)
print(f"Model Performance Metrics (Real Scale):")
print(f"R-squared (R2) Score: {r2:.4f}")
print(f"Explained Variance Score: {evs:.4f}")
print("-" * 50)

# Interpretation
if r2 > 0.7:
    print("The model explains a significant portion of the energy consumption variance.")
else:
    print("The model shows moderate predictive power; consider further feature engineering.")

```

---

Model Performance Metrics (Real Scale):  
R-squared (R2) Score: 0.5829  
Explained Variance Score: 0.5829

---

The model shows moderate predictive power; consider further feature engineering.

---

## From Determinism to Probability: Why Probabilistic Forecasting?

The exploratory analysis and initial regression tests highlighted that energy consumption is not a purely stochastic or isolated phenomenon, but a dynamic process with strong temporal dependencies. Framing the problem as a **Probabilistic Forecasting** task instead of classic **Regression** offers three crucial competitive advantages for a *Digital Twin*:

### 1. Modeling Temporal Dependency (Sequentiality)

Unlike standard regression, which treats each record as independent ( $P(y|x)$ ), forecasting recognizes that consumption at time  $t$  is intrinsically linked to previous instances ( $P(y_t|x_t, y_{t-1}, \dots, y_{t-n})$ ). The use of **LSTM** architectures allows mapping the building's thermal inertia and operational cycles, which a static regressor perceives as noise, artificially lowering the explained variance ( $R^2$ ).

### 2. Managing Uncertainty (Confidence Intervals)

In a critical monitoring system, knowing the point value (e.g., 75 kWh) is less useful than knowing the confidence interval. **Probabilistic Forecasting** provides a probability distribution: \* **Regression**: “I predict 75 kWh.” \* **Probabilistic**: “I predict 75 kWh with a standard deviation of 2 kWh.” This allows defining dynamic alert thresholds: a consumption of 80 kWh might be normal on a hot afternoon (high uncertainty) but critical during a winter night (low uncertainty).

### 3. Anomaly Detection Based on Likelihood

By transforming the task into a probabilistic one, fault detection evolves. We no longer look for a fixed percentage deviation but calculate the **Log-Likelihood** of the observed data. If the actual consumption falls in the “tails” of the distribution predicted by the model, the system flags an anomaly with a degree of statistical certainty, drastically reducing the false positives typical of threshold-based systems.

```
from sklearn.preprocessing import RobustScaler
from scipy import stats
```

```

# Loading and cleaning
df_nexus = df.copy()
df_nexus['Timestamp'] = pd.to_datetime(df_nexus['Timestamp'])

# Removing Outliers (Z-score > 3)
df_nexus = df_nexus[np.abs(stats.zscore(df_nexus['EnergyConsumption'])) < 3]

# Categorical Mapping
usage_map = {'On': 1, 'Off': 0}
df_nexus['HVACUsage'] = df_nexus['HVACUsage'].map(usage_map)
df_nexus['LightingUsage'] = df_nexus['LightingUsage'].map(usage_map)

# Feature Engineering: Cyclicality and Short-term Lag
df_nexus['hour_sin'] = np.sin(2 * np.pi * df_nexus['Timestamp'].dt.hour / 24)
df_nexus['hour_cos'] = np.cos(2 * np.pi * df_nexus['Timestamp'].dt.hour / 24)
df_nexus['Lag_1h'] = df_nexus['EnergyConsumption'].shift(1)

# Light smoothing to remove sensor noise
df_nexus['EnergyConsumption'] = df_nexus['EnergyConsumption'].rolling(window=2, center=True).mean()
df_nexus = df_nexus.dropna()

# Final Feature Selection
df_nexus['Temp_Trend'] = df_nexus['Temperature'].diff().fillna(0)
features = ['Temperature', 'Temp_Trend', 'Occupancy', 'HVACUsage', 'Lag_1h', 'hour_sin', 'hour_cos']
target = 'EnergyConsumption'

train_size = int(len(df_nexus) * 0.8)
train_df = df_nexus.iloc[:train_size]
test_df = df_nexus.iloc[train_size:]

scaler_X = RobustScaler()
scaler_y = RobustScaler()

X_train_raw = scaler_X.fit_transform(train_df[features])
y_train_raw = scaler_y.fit_transform(train_df[[target]])
X_test_raw = scaler_X.transform(test_df[features])
y_test_raw = scaler_y.transform(test_df[[target]])

WINDOW_SIZE = 12

def prepare_sequences(X, y, window):
    Xs, ys = [], []
    for i in range(len(X) - window):
        Xs.append(X[i : i + window])
        ys.append(y[i + window])

```

```

    return torch.tensor(np.array(Xs), dtype=torch.float32), \
           torch.tensor(np.array(ys), dtype=torch.float32)

X_train, y_train = prepare_sequences(X_train_raw, y_train_raw, WINDOW_SIZE)
X_test, y_test = prepare_sequences(X_test_raw, y_test_raw, WINDOW_SIZE)

train_loader = DataLoader(TensorDataset(X_train, y_train), batch_size=32, shuffle=True)

class NexusLSTM(nn.Module):
    def __init__(self, input_dim, hidden_dim=8):
        super(NexusLSTM, self).__init__()
        self.lstm = nn.LSTM(input_dim, hidden_dim, batch_first=True)
        self.bn = nn.BatchNorm1d(hidden_dim)
        self.dropout = nn.Dropout(0.4)
        self.fc = nn.Linear(hidden_dim, 1)

    def forward(self, x):
        _, (hn, _) = self.lstm(x)
        # hn[-1] is the last hidden state: [batch, hidden_dim]
        out = self.bn(hn[-1])
        out = self.dropout(out)
        return self.fc(out)

model = NexusLSTM(input_dim=len(features))
criterion = nn.HuberLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.5, pa

epochs = 200
history = {'train_loss': [], 'val_loss': []}

for epoch in range(epochs):
    model.train()
    train_running_loss = 0.0
    for batch_X, batch_y in train_loader:
        optimizer.zero_grad()
        preds = model(batch_X)
        loss = criterion(preds, batch_y)
        loss.backward()
        optimizer.step()
        train_running_loss += loss.item()

    avg_train_loss = train_running_loss / len(train_loader)

    model.eval()
    with torch.no_grad():
        val_preds = model(X_test)

```

```

        avg_val_loss = criterion(val_preds, y_test).item()
        scheduler.step(avg_val_loss)

    history['train_loss'].append(avg_train_loss)
    history['val_loss'].append(avg_val_loss)

    if (epoch + 1) % 50 == 0:
        print(f"Epoch [{epoch+1}/{epoch+1}] | Train: {avg_train_loss:.4f} | Val: {avg_val_loss:.4f}")

# Plot convergence
plt.figure(figsize=(10, 4))
plt.plot(history['train_loss'], label='Train Loss')
plt.plot(history['val_loss'], label='Validation Loss')
plt.title('Training vs Validation Loss')
plt.xlabel('Epochs')
plt.legend()
plt.show()

Epoch [50/200] | Train: 0.1856 | Val: 0.1889 | LR: 0.000125
Epoch [100/200] | Train: 0.1868 | Val: 0.1877 | LR: 0.000004
Epoch [150/200] | Train: 0.1846 | Val: 0.1881 | LR: 0.000000
Epoch [200/200] | Train: 0.1786 | Val: 0.1872 | LR: 0.000000

```




Figure 12: png

```

from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

# Set the model to evaluation mode
model.eval()

with torch.no_grad():

```

```

# Generate predictions using the test set
predictions_scaled = model(X_test)
predictions_numpy = predictions_scaled.cpu().numpy()
actuals_numpy = y_test.cpu().numpy()

# Inverse transform the scaled predictions and actual values to their original scale
predictions_kwh = scaler_y.inverse_transform(predictions_numpy)
actuals_kwh = scaler_y.inverse_transform(actuals_numpy)

# Plot the actual vs predicted energy consumption
plt.figure(figsize=(10, 6))
plt.scatter(actuals_kwh, predictions_kwh, alpha=0.5, color='teal', label='LSTM Predictions')
plt.plot([actuals_kwh.min(), actuals_kwh.max()], [actuals_kwh.min(), actuals_kwh.max()],
         'r--', lw=2, label='Perfect Prediction')

plt.xlabel('Actual Consumption (kWh)')
plt.ylabel('Predicted Consumption (kWh)')
plt.title('Digital Twin Analysis: Actual vs Predicted Energy (Real Scale)')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.6)
plt.show()

# Calculate Metrics
mae_kwh = mean_absolute_error(actuals_kwh, predictions_kwh)
mse_kwh = mean_squared_error(actuals_kwh, predictions_kwh)
r2 = r2_score(actuals_kwh, predictions_kwh)

print(f"--- Final Metrics of the Recurrent Model ---")
print(f"Mean Absolute Error (MAE): {mae_kwh:.2f} kWh")
print(f"Mean Squared Error (MSE): {mse_kwh:.2f} kWh^2")
print(f"R2 Score: {r2:.4f}")

--- Final Metrics of the Recurrent Model ---
Mean Absolute Error (MAE): 4.10 kWh
Mean Squared Error (MSE): 24.96 kWh^2
R2 Score: 0.2434

```

---

## Probabilistic Evolution: Transition to Negative Log Likelihood (NLL)

So far, we have treated energy consumption prediction as a **deterministic** problem: the model receives inputs and returns a single point value (Point Estimation). However, a building's energy consumption is an inherently **stochastic** phenomenon, influenced by unobserved variables and sensor noise.




Figure 13: png

The Negative Log Likelihood (NLL) loss function allows us to embrace this uncertainty by predicting a **probability distribution** for each time step instead of a single value. It is defined as follows for a Gaussian distribution:

$$NLL = \frac{1}{2} \log(2\pi\sigma^2) + \frac{(y - \mu)^2}{2\sigma^2}$$

Where: \*  $y$  = Actual consumption \*  $\mu$  = Predicted mean consumption \*  $\sigma^2$  = Predicted variance (uncertainty)

### Why NLL?

The **Gaussian Negative Log Likelihood** allows us to model predictions not as a single number, but as a **probability distribution**.

- **Instead of:**  $Output = y$
- **Now:**  $Output = \mathcal{N}(\mu, \sigma^2)$

The model will learn to provide two pieces of information for each time step:  
 1. **Mean ( $\mu$ ):** The most probable consumption value (the one we will use for MAE).  
 2. **Variance ( $\sigma^2$ ):** The model's degree of uncertainty. A high variance indicates that the data at that moment are chaotic or contradictory; a low variance indicates high confidence in the prediction.

## Benefits for the Digital Twin

- **Noise Resilience:** NLL penalizes errors less when the model has predicted high uncertainty ( $\sigma$ ), preventing network weights from being distorted by sudden outliers.
- **Confidence Intervals:** We can generate confidence bands (e.g.,  $\pm 2\sigma$ ) that indicate when real consumption is deviating from expected behavior (Anomaly Detection).

---

```
class NexusProbabilisticLSTM(nn.Module):
    def __init__(self, input_dim: int, hidden_dim: int = 12):
        super(NexusProbabilisticLSTM, self).__init__()
        self.lstm = nn.LSTM(input_dim, hidden_dim, batch_first=True)
        self.bn = nn.BatchNorm1d(hidden_dim)
        self.dropout = nn.Dropout(0.4)
        self.fc_mu = nn.Linear(hidden_dim, 1)
        self.fc_sigma = nn.Linear(hidden_dim, 1)
        self.softplus = nn.Softplus()

    def forward(self, x):
        _, (hn, _) = self.lstm(x)
        x = self.bn(hn[-1])
        x = self.dropout(x)
        mu = self.fc_mu(x)
        sigma = self.softplus(self.fc_sigma(x)) + 1e-6
        return mu, sigma

model = NexusProbabilisticLSTM(input_dim=len(features))
optimizer = torch.optim.Adam(model.parameters(), lr=0.0001, weight_decay=1e-4)
criterion = nn.GaussianNLLoss()
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.5, pa

# --- Training Loop with Validation and Scheduling ---
epochs = 250
history = {'train_loss': [], 'val_loss': []}

for epoch in range(epochs):
    model.train()
    train_running_loss = 0.0

    for batch_X, batch_y in train_loader:
        optimizer.zero_grad()
        mu, sigma = model(batch_X)
        # NLL requires mean, target, and variance (sigma^2)
        loss = criterion(mu, batch_y, sigma**2)
```

```

        loss.backward()

        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        optimizer.step()
        train_running_loss += loss.item()

        optimizer.step()
        train_running_loss += loss.item()

    # Validation phase
    model.eval()
    with torch.no_grad():
        val_mu, val_sigma = model(X_test)
        val_loss = criterion(val_mu, y_test, val_sigma**2).item()

    # Update scheduler based on validation performance
    scheduler.step(val_loss)

    avg_train_loss = train_running_loss / len(train_loader)
    history['train_loss'].append(avg_train_loss)
    history['val_loss'].append(val_loss)

    if (epoch + 1) % 20 == 0:
        curr_lr = optimizer.param_groups[0]['lr']
        print(f"Epoch [{epoch+1}/{epochs}] | Train: {avg_train_loss:.4f} | Val: {val_loss:.4f}")

    Epoch [20/250] | Train: 0.4457 | Val: 0.0718 | LR: 0.000100
    Epoch [40/250] | Train: 0.2482 | Val: 0.0358 | LR: 0.000100
    Epoch [60/250] | Train: 0.0928 | Val: 0.0232 | LR: 0.000100
    Epoch [80/250] | Train: 0.1085 | Val: 0.0170 | LR: 0.000100
    Epoch [100/250] | Train: -0.0062 | Val: 0.0229 | LR: 0.000050
    Epoch [120/250] | Train: 0.0031 | Val: 0.0267 | LR: 0.000050
    Epoch [140/250] | Train: -0.0128 | Val: 0.0221 | LR: 0.000025
    Epoch [160/250] | Train: -0.0084 | Val: 0.0216 | LR: 0.000013
    Epoch [180/250] | Train: 0.0103 | Val: 0.0281 | LR: 0.000006
    Epoch [200/250] | Train: 0.0251 | Val: 0.0209 | LR: 0.000003
    Epoch [220/250] | Train: -0.0168 | Val: 0.0255 | LR: 0.000003
    Epoch [240/250] | Train: -0.0100 | Val: 0.0228 | LR: 0.000002

    model.eval()
    with torch.no_grad():
        mu_scaled, sigma_scaled = model(X_test)
        predictions_kwh = scaler_y.inverse_transform(mu_scaled.cpu().numpy())
        actuals_kwh = scaler_y.inverse_transform(y_test.cpu().numpy())
        # Scale uncertainty back to kWh
        std_kwh = sigma_scaled.cpu().numpy() * scaler_y.scale_

```

```

mae_final = mean_absolute_error(actuals_kwh, predictions_kwh)
r2_final = r2_score(actuals_kwh, predictions_kwh)

print(f"\n--- Final Probabilistic Results with Scheduling ---")
print(f"MAE (Mean): {mae_final:.2f} kWh")
print(f"R2 Score: {r2_final:.4f}")

# Plotting the results
plt.figure(figsize=(15, 6))
plt.plot(actuals_kwh[:100], label='Actual', color='black', alpha=0.4)
plt.plot(predictions_kwh[:100], label='Prediction', color='teal', linewidth=2)
plt.fill_between(
    range(100),
    (predictions_kwh[:100] - 1.96 * std_kwh[:100]).flatten(),
    (predictions_kwh[:100] + 1.96 * std_kwh[:100]).flatten(),
    color='teal', alpha=0.2, label='95% Confidence Interval'
)
plt.title('Nexus Digital Twin: Probabilistic NLL')
plt.legend()
plt.show()

--- Final Probabilistic Results with Scheduling ---
MAE (Mean): 4.09 kWh
R2 Score: 0.2495

```




Figure 14: png

```

# Analysis of Residuals by Hour
test_results = test_df.iloc[WINDOW_SIZE:].copy()
test_results['Error'] = np.abs(actuals_kwh - predictions_kwh)
test_results['Hour'] = test_results['Timestamp'].dt.hour

# Calculate average MAE per hour

```

```

hourly_mae = test_results.groupby('Hour')['Error'].mean()

plt.figure(figsize=(12, 5))
hourly_mae.plot(kind='bar', color='teal', alpha=0.7)
plt.axhline(y=mae_final, color='r', linestyle='--', label='Global MAE')
plt.title('MAE Distribution by Hour of the Day')
plt.ylabel('MAE (kWh)')
plt.xlabel('Hour of Day')
plt.legend()
plt.grid(axis='y', alpha=0.3)
plt.show()

```




Figure 15: png

### Hourly Performance Analysis

The bar chart highlighting the **MAE by Hour of Day** reveals a non-uniform distribution of uncertainty:

- \* **Stability during Low Activity:** The model is most accurate during off-peak hours (e.g., 02:00, 14:00, 23:00), where the error drops below 3 kWh. This indicates that the base thermal load of the building is well understood.
- \* **Transition Noise (07:00 - 13:00):** A sustained error above the global average during the morning suggests that the “ramp-up” phase of the building is stochastic. Human arrivals and HVAC ignition timing vary slightly each day, creating a ceiling of unpredictability.
- \* **Peak Anomaly (20:00):** The significant spike at 8 PM suggests an unobserved behavioral variable (e.g., specific evening cleaning shifts or security protocols) that the current feature set cannot account for.

```

# Residual Distribution Check
plt.figure(figsize=(10, 5))
residuals = actuals_kwh - predictions_kwh

```

```

plt.hist(residuals, bins=50, color='orange', edgecolor='black', alpha=0.7)
plt.title('Residuals Distribution (Actual - Predicted)')
plt.xlabel('Error (kWh)')
plt.ylabel('Frequency')
plt.show()

```




Figure 16: png

### Residuals Distribution and Multimodality

The histogram of residuals (**Actual - Predicted**) provides the “statistical signature” of our model:

- \* **Zero-Centered Bias:** The distribution is centered at zero, proving that the model is **unbiased**. It does not systematically overestimate or underestimate consumption.
- \* **Multimodal Nature:** Instead of a perfect Gaussian curve, we see multiple peaks (e.g., at -3, +2, and +6 kWh). This confirms that the dataset contains several “operational modes” (e.g., different types of occupancy days) that overlap.
- \* **Positive Skewness (Underestimation):** The longer tail on the right indicates that the model is conservative; it tends to underpredict extreme energy spikes rather than overpredicting them. This is a typical effect of the **Dropout** and **Huber Loss** regularization used to maintain stability.